

Operating System Process and Resource Management

Inter-Process Communication



Course Objectives

- ✓ Understand the problem of concurrent access
- ✓ Understand the resolution algorithms
- ✓ Find optimized solutions



Course Plan

1. Locking and Deadlocking
2. Remove Active Waiting



1. Locking and Deadlocking



1. Locking and Deadlocking

Introduction

- Modeling real systems in a sequential program must:
 - Define in its code the management of the components of its environment
 - Be able to modify the program when new activities are included
 - Be able to modify the program when changing the architecture (number of CPU, *etc.*)
- The real world is made of competing programs that work together, for example:
 - A switch uses information from multiple sensors simultaneously and provides orders to multiple actors
 - A reservation system considers that there are several agencies in parallel
- Synchronizing processes can lead to problems such as **locking**, **deadlocking** or **starvation**

1. Locking and Deadlocking

Introduction

- Our current programs are multi-tasking and inter-tasking, they can:
 - Share and exchange information
 - Synchronize to respect the use and protection of data
- We have 2 models of systems:
 - Centralized, that communicate and synchronize using a common memory
 - Distributed, that communicate and synchronize by sending messages
- **Multi-tasking** = on a single processor
- **Parallel** = several processors with a common memory
- **Distributed** = remote machines (networked) and without a common memory

1. Locking and Deadlocking

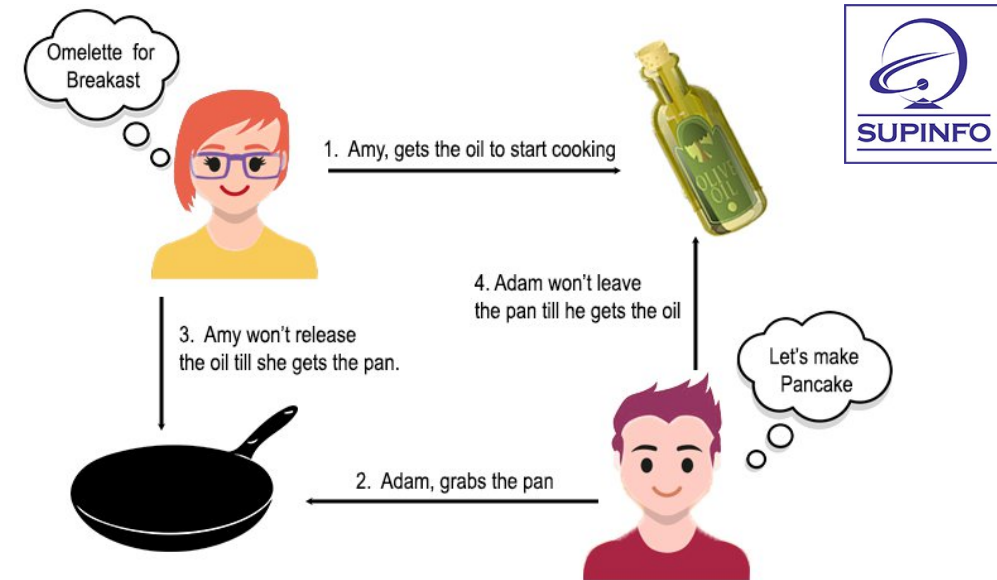
Examples

User	Action	Output
Bob	getState(room1)	free
Alice	getState(room1)	free
Bob	book(room1)	OK
Alice	book(room1)	OK

} 2 reservations for the same room!

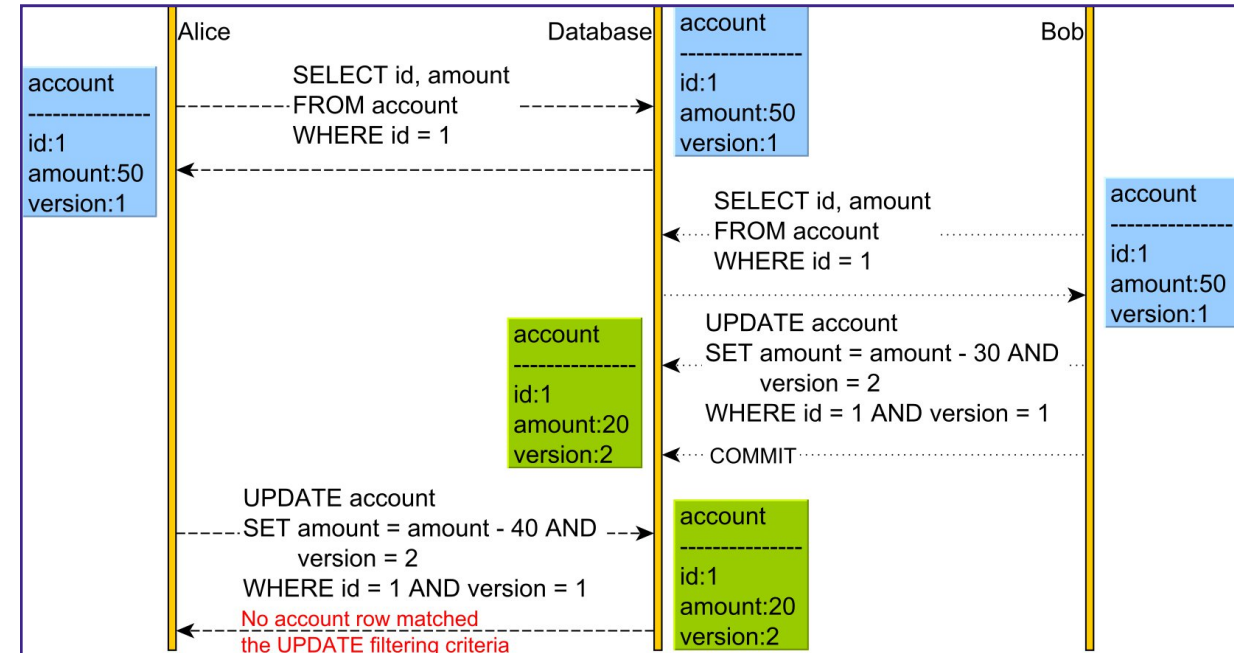
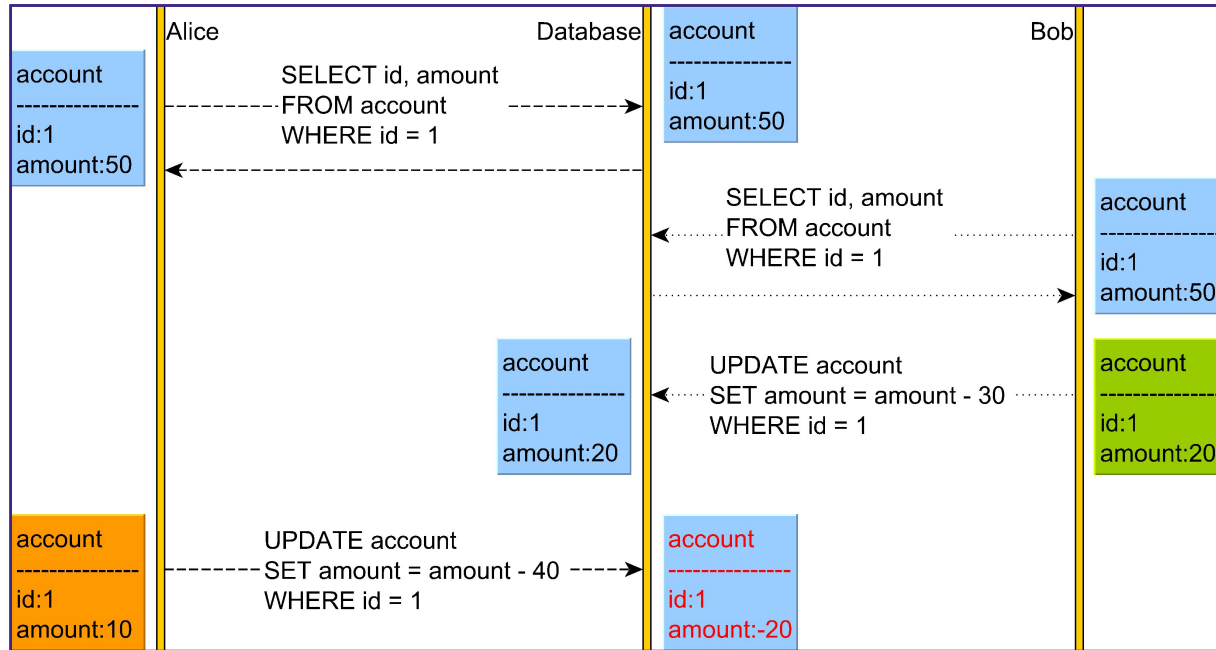
User	Action	Output
Bob	getState(room1)	free
Alice	getState(room1)	free
Bob	book(room1)	OK
Alice	book(room1)	KO

} Updated state or back check



1. Locking and Deadlocking

Examples



1. Locking and Deadlocking

Critical section

- The previous examples show that data protection systems are needed, we denote several (like for centralized systems):
 - **Test and Set**, low-level (processor)
 - **Interrupt masking**, dangerous (kernel)
 - **Semaphore**, common and simple
 - **Monitor**, high-level (language level)
- All these systems act on data that must be protected (program level)
- The **critical section** (CS) is the part of a program whose execution should not **intermingle** with other programs
 - Once a task enters a critical section, it must be allowed to complete that section without allowing other tasks to execute the same data

1. Locking and Deadlocking

Problem generalization

- When a process manipulates a sensitive data, we say that it is in a **critical section**
- When a process does not manipulate sensitive data and it goes out of the critical section, we say that it is in the **remainder section** (RS)

Repeat

- Entry section
- **Critical section**
- Exit section
- **Remainder section**

Forever

Critical section

1. Bob requests a room reservation

Entry section

2. Database states that room1 is available

3. Room1 is allocated to Bob and tagged as reserved

Exit section

1. Locking and Deadlocking

Relevance

- We find this problem in:
 - Databases
 - Resource sharing (files, connections, networks, *etc.*)
 - Automats (like ATM)
 - Hardware (hard disk, *etc.*)
 - Software (MMORPG, client/server, *etc.*)
 - *etc.*



1. Locking and Deadlocking

Necessary criteria for a valid solution

- **Mutual exclusion:** at anytime, at most one process can be in a critical section for a given variable
- **Non-interference:** if a process stops in its remainder section, this should not affect the other processes
- The problem in the critical section is to find a mutual exclusion algorithm
 - Software-provided solutions
 - Hardware-provided solutions
 - OS-provided solutions

1. Locking and Deadlocking


First algorithmic solution

- We start with 2 processes: P0 and P1
- When we consider the task P_i , P_j will always refer to another task ($i \neq j$)
- This software solution will be algorithmically composed of processes giving each other the turn:
 1. The shared variable token is initialized to 0 or 1
 2. The CS of P_i is executed if and only if token = i
 3. P_i is waiting if P_j is in CS

```
Repeat # $P_i$  process
    While(token  $\neq i$ ) { #wait }
    CS
    token = j
    RS
Forever
```

1. Locking and Deadlocking

First algorithmic solution



1. Locking and Deadlocking

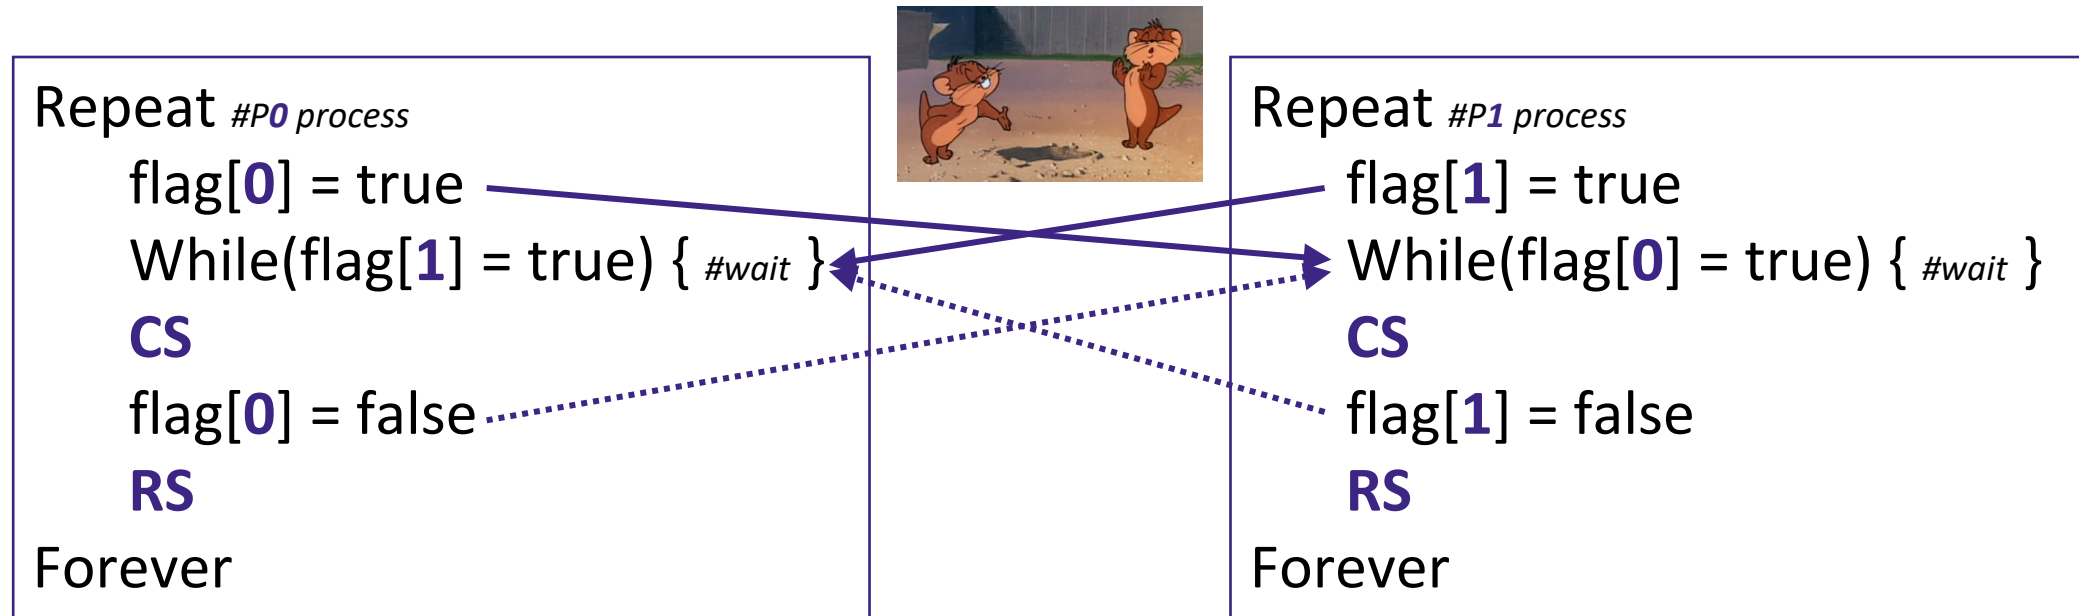
Courtesy excess

- Generalization to n processes: each time, before a process can enter its critical section, it must wait until all the others have had this opportunity
 - A Boolean variable per process: ***flag[0]*** and ***flag[1]***
 - P_i signals that it wants to execute its CS: ***flag[i] = true***

```
Repeat # $P_i$  process
    flag[i] = true
    While(flag[j] = true) { #wait }
    CS
    flag[i] = false
    RS
Forever
```

1. Locking and Deadlocking

Courtesy excess



1. Locking and Deadlocking

Dekker algorithm

- Initialization: c_0 , c_1 , token = 0 (previous locking broken thanks to the token)

Repeat *#P₀ process*

NCS0: *#non-critical section*

$c_0 = \text{false}$

While($c_1 = \text{false}$) {

$c_0 = \text{true}$

While(token = 1) { *#wait* }

$c_0 = \text{false}$

}

CS0: *#critical section*

token = 1

$c_0 = \text{true}$

Forever

Repeat *#P₁ process*

NCS1: *#non-critical section*

$c_1 = \text{false}$

While($c_0 = \text{false}$) {

$c_1 = \text{true}$

While(token = 0) { *#wait* }

$c_1 = \text{false}$

}

CS1: *#critical section*

token = 0

$c_1 = \text{true}$

Forever

1. Locking and Deadlocking

Dekker algorithm

- The correct operation of Dekker's algorithm depends on the fact that each process always ends up progressing (fairness)
- A simple fairness assumption, sufficient to reason about the Dekker algorithm, is as follows: any process that has the ability to execute an instruction will always end up doing so
- Among the properties we specify for a concurrent program, we distinguish the following 2 categories:
 - **Safety properties**: these are properties that express that certain undesirable states are never reached, they do not depend on the use of a fairness assumption
 - **Liveliness properties**: these express that desirable states will inevitably be reached, they are usually only true in the presence of a fairness assumption

1. Locking and Deadlocking

Test-and-set instruction

- An algorithm using test-and-set for mutual exclusion:
 - Shared variable **i** is initialized to 0
 - It is the first **P_i** (which sets **i** to 1) that enters CS

Non-divisible and
unbreakable atomic
instruction

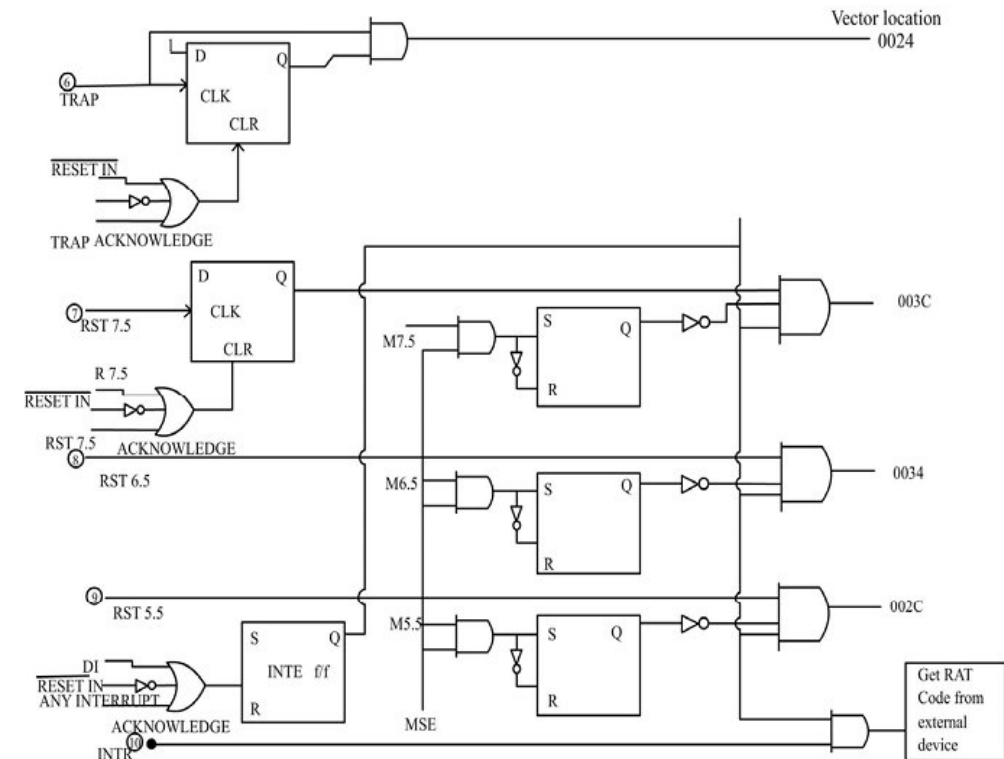
```
bool TestAndSet(int &i)
{
    If(i = 0) {
        i = 1
        return true
    } else {
        return false
    }
}
```

```
Repeat #Pi process
    While(TestAndSet(i) = false) { #wait }
    CS
    i = 0
    RS
Forever
```

1. Locking and Deadlocking

Need to use other methods

- The major drawback of software solutions is **active waiting**
- Hardware solutions are for mutual exclusion
- **Interrupt masking**: a process that wants to enter a critical section inhibits interrupts to conserve the processor until it restores them



1. Locking and Deadlocking

Exercise

- Define uniprogramming and multiprogramming
- Why is a process table necessary in a time-sharing system?
- Is a process table also required in a personal system where only one process exists, with access to the entire machine during its execution?



1. Locking and Deadlocking

Questions



2. Remove Active Waiting



2. Remove Active Waiting

Semaphore

- Semaphore was proposed by Dijkstra in 1965 which is a very significant technique to manage concurrent processes, it is an elementary structure of synchronization (limiting active waiting)
- A semaphore S is a set containing:
 - An internal (integer) counter: $c(S)$
 - A process queue : $q(S)$
 - 2 atomic operations/primitives (non-breakable execution): $P(S)$ and $V(S)$
 - An atomic operation of initialization (and creation) of the counter
- A semaphore S is said to be binary if its counter remains lower or equal to 1

2. Remove Active Waiting

Semaphore

- **P(S)** operation is also called wait, sleep, or down operation, and **V(S)** operation is also called signal, wake-up, or up operation
- Let p be the process that executes **P(S)** or **V(S)** with **c(S)** arbitrarily initialized:

P(S)

```
c(S) = c(S)-1
If(c(S)<0) {
    state(p) = locked
    input(p, q(S))
}
```

V(S)

```
c(S) = c(S)+1
If(c(S)≤0) {
    output(p, q(S))
    state(p) = eligible
    input(p, q(eligible))
}
```

2. Remove Active Waiting

Semaphore

- Binary (**mutex**) semaphore:
 - The purpose of exclusion for each process is to protect access to a single resource (memory, disk, *etc.*)
 - The counter **c(S)** is initialized to 1
- Counting (**synchronization**) semaphore:
 - The purpose of synchronization is that one process must wait for another to continue (or start) its execution
 - The counter **c(S)** is initialized to 0

$P(S)$ *#Takes the critical section*
CS
 $V(S)$ *#Releases the critical section*

P0

job**0**()
 $V(S)$ *#Wakes up P1*

P1

$P(S)$ *#Waits P0*
job**1**()

2. Remove Active Waiting

Rendezvous

- The goal of the rendezvous (generalization of the previous problem) is that a process waits for n other processes to be at a specific execution to continue its execution:
 - **Sync**, a semaphore initialized to 0: $Sync = S(0)$
 - **Mutex**, a semaphore initialized to 1: $Mutex = S(1)$
 - **Wait**, a semaphore initialized to 0: $Wait = S(0)$
 - A counter (integer) **nb** initialized to 0: $nb = 0$
 - An indicator (integer) the **n** processes waiting: n

P_i

```
P(Mutex)
nb = nb+1
If(nb = n) {
    V(Sync)
}
nb = 0
V(Mutex)
P(Wait)
```

$Prdv$

```
...
P(Sync)
V(Wait, n)
...
```

2. Remove Active Waiting

Deadlocking risk

- For the deadlocks we have:
 - S0 and S1, 2 mutex semaphores
 - **P()** always executed in the same order

P_i

P(S0)
P(S1)
CS
V(S1)
V(S0)



P_j

P(S1)
P(S0)
CS
V(S0)
V(S1)

2. Remove Active Waiting

Deadlocking risk


- To prevent and respond to deadlocking, the calls between the processes must be represented:



- If there is a call from one process to another and vice versa, then there is a deadlock
- To limit this locking, we must add the operations **V()** to release missing resources

2. Remove Active Waiting

Producer-consumer problem



- 1-slot model:
 - The producer and the consumer are 2 cyclic processes

Producer

```
produce(messageP)  
deliver(slot, messageP)
```

Consumer

```
remove(slot,  
messageC)  
consume(messageC)
```

2. Remove Active Waiting

Producer-consumer problem

- Synchronized producer-consumer model with a shared array (deliver/remove) with 1 slot; we use 2 semaphores, **full** and **empty**, initialized to 0 and 1 (**empty** indicates if the slot of the shared array is empty, inversely for **full**)

Producer

```
P(empty)  
produce(messageP)  
deliver(slot, messageP)  
V(full)
```

or

```
produce(messageP)  
P(empty)  
deliver(slot, messageP)  
V(full)
```

Consumer

```
P(full)  
remove(slot,  
messageC)  
consume(messageC)
```

V(**empty**) or

```
P(full)  
remove(slot,  
messageC)  
V(empty)
```

consume(messageC)

2. Remove Active Waiting

Monitor

- The monitor is an advanced synchronization tool, introduced by Brinch Hansen in 1972-73 and Hoare in 1974; it consists of :
 - State variables
 - Internal procedures
 - External procedures (entry points)
 - Conditions
 - Synchronization primitives
- An example of a monitor for incrementing and decrementing a variable *i*

Monitor incr_decr

```
incr_decr: monitor
  var i: integer

  procedure increment
  begin
    i = i+1
  end

  procedure decrement
  begin
    i = i-1
  end

  begin
    i = 0
  end
end incr_decr
```


2. Remove Active Waiting

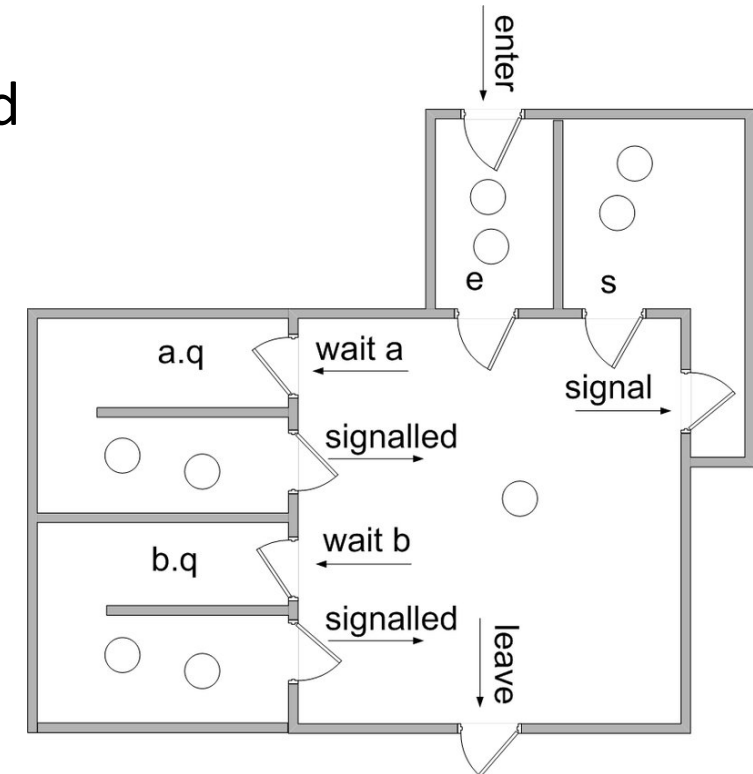
Monitor

- Each monitor procedure is executed in mutual exclusion
- The access to the monitor is done in mutual exclusion, except if a process executes the **wait** primitive
- Special monitor instructions: the **empty**, **wait** and **signal** primitives on condition type variables
- Consider **c** as a condition:
 - **c.wait**: lock the process and put it in waiting for the condition **c**
 - **c.empty**: *true* if no process is waiting for condition **c**, *false* otherwise
 - **c.signal**: if not **c.empty** then wake up a process waiting for condition **c**

2. Remove Active Waiting

Monitor

- Inside the monitors each monitor has a global waiting queue and each variable condition **c** acts on a waiting queue:
 - **c.wait**: put the process in the waiting queue associated with condition **c**
 - **c.empty**: check if the waiting queue associated with condition **c** is empty
 - **c.signal**: take the next process out of the queue associated with the condition **c**



2. Remove Active Waiting

Monitor

- Only one process is active at a time within a monitor, so it is not autonomous to implement the **signal** primitive
- A monitor is usually implemented with semaphores
- Such a system is thus built by successive levels of abstraction, and a level i is implemented by the level $i-1$ primitives



2. Remove Active Waiting

Exercise

- Write a monitor to perform the rendezvous of n processes: the purpose of this monitor is to allow the waiting of n processes according to the condition of the number of processes waiting and the processes available
- Solve the producer-consumer problem with monitors using a single slot




2. Remove Active Waiting

Questions



Operating System Process and Resource Management

Inter-Process Communication



Thank you for your attention